

# Curs 6

Programare Paralela si Distribuita

Variabile ThreadLocal  
Intreruperi  
Forme de sincronizare  
Bariera de sincronizare  
Rendez-vous

Exchanger

# ThreadLocal

- Variabile locale fiecarui thread (pt fiecare thread se creeaza o copie)
- Instantele sunt in general campuri `private static` in clase care doresc sa asocieze cate o stare fiecarui thread
- Trei metode publice : `set`, `get`, `remove`

```
private ThreadLocal<String> myThreadLocal = new ThreadLocal<String>();
```

```
myThreadLocal.set("Hello ThreadLocal");
```

```
String threadLocalValue = myThreadLocal.get();
```

- Valorile obiectului de tip `ThreadLocal` sunt vizibile doar de catre thread-ul care seteaza valoarea.
- Setare valoare initiala:

```
private ThreadLocal myThreadLocal = new ThreadLocal<String>() {  
    @Override protected String initialValue() {  
        return "This is the initial value";  
    }  
};
```

# Exemplul – Atomic + ThreadLocal

```
import java.util.concurrent.atomic.AtomicInteger;

public class ThreadId {
    // Atomic integer containing the next thread ID to be assigned
    private static final AtomicInteger nextId = new AtomicInteger(0);

    // Thread local variable
    private static final ThreadLocal<Integer> threadId = new ThreadLocal<Integer> () {
        @Override
        protected Integer initialValue() {
            return nextId.getAndIncrement();
        }
    };

    // Returns the current thread's unique ID, assigning it if necessary
    public static int get() {
        return threadId.get();
    }
}
```

# Intreruperi

- O intrerupere ( *interrupt* ) este o indicatie pentru un thread ca ar trebui sa se opreasca si ... sa faca altceva (de ex. sa se termine).

`public void interrupt()`

`public static boolean interrupted()`

`public boolean isInterrupted()`

*“There is no way in Java to terminate a thread unless the thread exits by itself.”*

# Intreruperi

- mecanismul de intrerupere foloseste un flag intern -> the *interrupt status*.
- Atunci cand se apeleaza `Thread.interrupt` se seteaza acest flag.
- Atunci cand se verifica intreruperea prin metoda statica `Thread.interrupted`, *<interrupt status>* este sters.
- Metoda nestatica `isInterrupted`, care este folosita de catre un thread pt a verifica statusul (*interrupt status*) al altuia nu schimba flagul.
- Prin conventie, orice metoda care se termina (exit) aruncand o exceptie de tip `InterruptedException` sterge "interrupt status".
- Totusi este posibil ca acesta sa fie imediat setat din nou de catre alt thread care invoca o metoda *interrupt*.

# Exemplu

```
public class SimpleThreads {

    static void threadMessage(String message) {
        String threadName = Thread.currentThread().getName();
        System.out.format("%s: %s %n", threadName, message);
    }

    private static class MessageLoop implements Runnable {
        public void run() {
            String importantInfo[] = {
                "Studentii sunt prezenti.",
                "Examenul este greu.",
                "Vacanta este asteptata.",
                "Exista concurenta."
            };
            try {
                for (int i = 0; i < importantInfo.length; i++) {
                    // Pause for 4 seconds
                    Thread.sleep(4000);
                    // Print a message
                    threadMessage(importantInfo[i]);
                }
            } catch (InterruptedException e) {
                threadMessage("I wasn't done!");
            }
        }
    }
}
```

```

public static void main(String args[])
throws InterruptedException {

    // Delay, in milliseconds before
    // we interrupt MessageLoop
    // thread (default one hour).

    long patience = 1000 * 60 * 60;

    threadMessage(
        "Starting MessageLoop thread");

    long startTime = System.currentTimeMillis();

    Thread t = new Thread(new MessageLoop());

    t.start();

    threadMessage(
        "Waiting for MessageLoop thread to finish");

```

```

while (t.isAlive()) {
    threadMessage("Still waiting...");

    t.join(1000);

    if ((
        (System.currentTimeMillis() - startTime) > patience)
        && t.isAlive())
    {
        threadMessage("Tired of waiting!");
        t.interrupt();
        // Shouldn't be long now
        // -- wait indefinitely

        t.join();
    }
    threadMessage("Finally!");
}

```

# Interactiuni: Waits, Notification, Interruption

- Notificarile nu se pot pierde din cauza intreruperilor.
- Presupunem ca un set de threaduri **s** este in **wait set** a lui **m**, si alt thread executa notificare pe **m**.

Atunci fie:

- cel putin un thread din **s** iese **normal din wait**, sau
- **toate threadurile din s ies din wait** aruncand exceptie **InterruptedException**
- Daca un thread este atat intrerupt cat si trezit prin notificare si el iese din **wait** aruncand o exceptie atunci un alt thread din **wait set** va fi notificat.
- Daca thread **t** a fost sters din **wait set** a lui **m** din cauza unei intreruperi atunci **interruption status** al lui **t** este setat la **false** si se iese din **wait** cu aruncarea unei exceptii de tip **InterruptedException**.



# Conditionare - reguli

- folosirea operatiilor `wait` doar in cicluri care se termina atunci cand anumite conditii logice sunt indeplinite!
- fiecare thread trebuie sa determine o ordine intre evenimentele care pot cauza ca el sa fie sters din *wait set*.

De exemplu: daca *t* este in *wait set* al obiectului ***o***, atunci cand apare atat o intrerupere a lui ***t*** cat si o notificare a lui ***o***, trebuie sa existe o ordine intre aceste evenimente.

- | Daca *interrupt* este considerata prima, atunci pana la urma *t* iese din `wait` aruncand `InterruptedException`, si alt thread din *wait set* a lui *o* (daca mai exista vreunul) va primi notificarea.
- | Daca sunt invers ordonate atunci *t* iese normal din `wait` si intreruperea este in asteptare (pana se verifica starea).

## Concluzii

Un thread  $t$  poate fi sters din *wait set* al unui obiect  $o$  ca urmare a uneia din urm. actiuni (si apoi isi va relua executia) :

- ◆ O actiune *notify* asupra lui  $o$  in care  $t$  este selectat spre a fi sters din *wait set*.
- ◆ O actiune *notifyAll* asupra lui  $o$  .
- ◆ O actiune de intrerupere realizata de  $t$ .
- ◆ Trecerea timpului (argument) specificat la apelul lui *wait*.
- ◆ Operatii "spurious wake-ups" (*no apparent reason*) –/rare.

([https://en.wikipedia.org/wiki/Spurious\\_wakeup](https://en.wikipedia.org/wiki/Spurious_wakeup))

## *Nested monitor lockout*

- Thread 1 blocheaza A
  - Thread 1 blocheaza B (in timp se ramane blocajul pe A)
    - Thread 1 decide sa astepte un semnal de la un alt thread
    - Thread 1 apeleaza `B.wait()` => elibereaza B dar nu si pe A.
- Thread 2 trebuie sa blocheze atat pe A cat si pe B (in ordine) pentru a trimite un semnal catre Thread 1.
- Thread 2 nu poate bloca pe A (pt ca Thread 1 are blocajul lui A).
- Thread 2 ramane blocat indefinit asteptand ca A sa fie eliberat.
- Thread 1 ramane blocat indefinit asteptand un semnal de la Thread 2, si astfel nu elibereaza pe A, etc.

# Diferenta intre Deadlock si *Nested monitor lockout*

- In cazul *deadlock*, 2 (sau mai multe) threaduri se asteapta unul pe altul sa elibereze blocajul.
- In cazul *nested monitor lockout*:
  - Thread 1 blocheaza A, si asteapta un semnal de la Thread 2.
  - Thread 2 asteapta sa-l blocheze pe A pentru a trimite semnalul catre Thread 1.

# Bariera de sincronizare

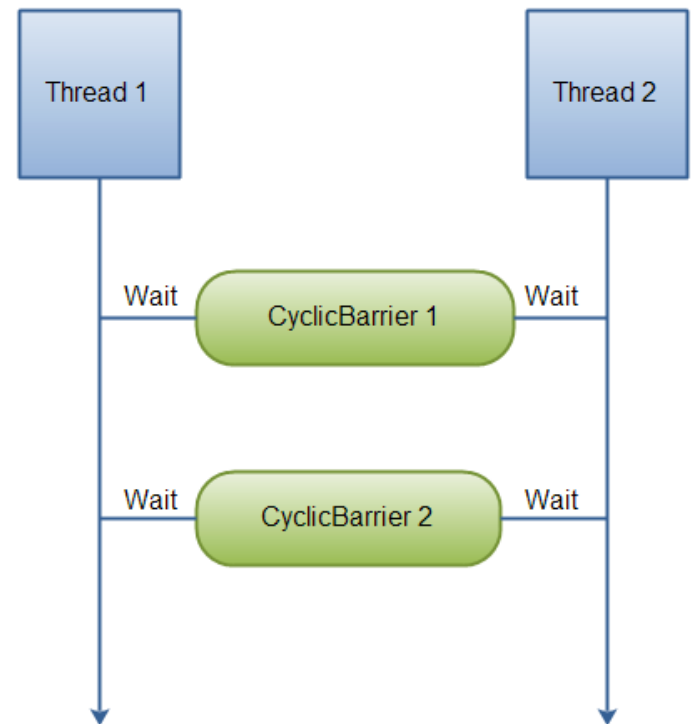
```
CyclicBarrier barrier = new CyclicBarrier(2);  
// 2 = no_of_threads_to_wait_for
```

```
barrier.await();
```

```
barrier.await(10, TimeUnit.SECONDS);
```

## Bariera de sincronizare:

- Bariera secventiala – pt implementare se fol. in general 2 variabile – {no\_threads(0..n), state(stop/pass)}
- Bariera ierarhica (tree-barrier)



Jacob Jenkov Tutorial

`java.util.concurrent>`

## `java.util.concurrent.CyclicBarrier` (java documentation)

- A synchronization aid that allows a set of threads to all wait for each other to reach a common barrier point. CyclicBarriers are useful in programs involving a fixed sized party of threads that must occasionally wait for each other.
- The barrier is called cyclic because it can be re-used after the waiting threads are released.
- A CyclicBarrier supports an optional Runnable command that is run once per barrier point, after the last thread in the party arrives, but before any threads are released. This barrier action is useful for updating shared-state before any of the parties continue.
  - `CyclicBarrier(int parties)`
  - `CyclicBarrier(int parties, Runnable barrierAction)`

# java.util.concurrent.CountDownLatch

## (java documentation)

- A synchronization aid that allows one or more threads to wait until a set of operations being performed in other threads completes.
- A `CountDownLatch` is initialized with a given count. The `await` methods block until the current `count` reaches zero due to invocations of the `countDown()` method, after which all waiting threads are released and any subsequent invocations of `await` return immediately.
  - This is a one-shot phenomenon -- the count cannot be reset. If you need a version that resets the count, consider using a `CyclicBarrier`.
- A `CountDownLatch` is a versatile synchronization tool and can be used for a number of purposes.
  - A `CountDownLatch` initialized with a count of one serves as a simple on/off latch, or gate: all threads invoking `await` wait at the gate until it is opened by a thread invoking `countDown()`.
  - A `CountDownLatch` initialized to `N` can be used to make one thread wait until `N` threads have completed some action, or some action has been completed `N` times.
- A useful property of a `CountDownLatch` is that it doesn't require that threads calling `countDown` wait for the count to reach zero before proceeding, it simply prevents any thread from proceeding past an `await` until all threads could pass.

# CountDownLatch (Java)

```
final CountDownLatch latch = new CountDownLatch(5);
```

```
// making two threads for 2 services
```

```
Thread serviceOneThread = new Thread(new ServiceOne(latch, 2));
```

```
Thread serviceTwoThread = new Thread(new ServiceTwo(latch, 3));
```

```
serviceOneThread.start();
```

```
serviceTwoThread.start();
```

```
// latch waits till the count becomes 0
```

```
// this way it can make sure that the execution of main thread only
```

```
// finishes once 2 services have executed
```

```
try {
```


```
    latch.await(); ...
```



# Conceptul de intalnire (rendez-vous)

- Conceptul de întâlnire (rendez-vous) a fost introdus initial in limbajul Ada pentru a facilita comunicarea între două task-uri.
  - a) Procesul B este gata să transmită informațiile, dar procesul A nu le-a cerut încă. In acest caz, procesul B rămâne în așteptare până când procesul A i le cere.
  - b) Procesul B este gata să transmită informațiile cerute, iar procesul A cere aceste date. In acest caz, se realizează un *rendez-vous*, cele două procese lucrează sincron până când își termină schimbul, după care fiecare își continuă activitatea independent.
  - c) Procesul A a lansat o cerere, dar procesul B nu este în măsură să-i furnizeze informațiile solicitate. In acest caz, A rămâne în așteptare până la întâlnirea cu B.

# Java Exchanger <-> Rendez-Vous



Jacob Jenkov Tutorial

Thread-0 exchanged A for B  
Thread-1 exchanged B for A

# Exchanger

## (java doc)

`public V exchange(V x)`      `throws InterruptedException`

- Waits for another thread to arrive at this exchange point (unless the current thread is interrupted), and then transfers the given object to it, receiving its object in return.
- If another thread is already waiting at the exchange point then it is resumed for thread scheduling purposes and receives the object passed in by the current thread.
  - The current thread returns immediately, receiving the object passed to the exchange by that other thread.

If no other thread is already waiting at the exchange then the current thread is disabled for thread scheduling purposes and lies dormant until one of two things happens:

- Some other thread enters the exchange; or
- Some other thread interrupts the current thread.

If the current thread:

- has its interrupted status set on entry to this method; or
- is interrupted while waiting for the exchange,
  - then `InterruptedException` is thrown and the current thread's interrupted status is cleared.

```
Exchanger exchanger = new Exchanger();
```

```
ExchangerRunnable exchangerRunnable1 = new ExchangerRunnable(exchanger, "A");
```

```
ExchangerRunnable exchangerRunnable2 = new ExchangerRunnable(exchanger, "B");
```

```
new Thread(exchangerRunnable1).start();
```

```
new Thread(exchangerRunnable2).start();
```

```

public class ExchangerRunnable implements Runnable{
    Exchanger exchanger = null;
    Object object = null;

    public ExchangerRunnable(Exchanger exchanger, Object object) {
        this.exchanger = exchanger;
        this.object = object;
    }

    public void run() {
        try {
            Object previous = this.object;
            this.object = this.exchanger.exchange (this.object);

            System.out.println(
                Thread.currentThread().getName() +
                " exchanged " + previous + " for " + this.object
            );
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}

```

# Class SynchronousQueue (->Rendez-vous)

## Java doc

[java.lang.Object](#)

[java.util.AbstractCollection<E>](#)

[java.util.AbstractQueue<E>](#)

java.util.concurrent.SynchronousQueue<E>

A blocking queue in which each insert operation must wait for a corresponding remove operation by another thread, and vice versa. **A synchronous queue does not have any internal capacity, not even a capacity of one.** You cannot peek at a synchronous queue because an element is only present when you try to remove it; **you cannot insert an element (using any method) unless another thread is trying to remove it**; you cannot iterate as there is nothing to iterate. The *head* of the queue is the element that the first queued inserting thread is trying to add to the queue; if there is no such queued thread then no element is available for removal and poll() will return null. For purposes of other Collection methods (for example contains), a SynchronousQueue acts as an empty collection. This queue does not permit null elements [<https://docs.oracle.com/javase/7/docs/api/java/util/concurrent/SynchronousQueue.html>]

- boolean [offer\(E e\)](#)
  - Inserts the specified element into this queue, if another thread is waiting to receive it.
- void [put\(E o\)](#)
  - Adds the specified element to this queue, waiting if necessary for another thread to receive it.
- [E poll\(\)](#)
  - Retrieves and removes the head of this queue, if another thread is currently making an element available.

# Exemplu

```
public class SynchronousQueueDemo{ public static void main(String args[]) {  
    final SynchronousQueue<String> queue = new SynchronousQueue<String>();  
    Thread producer = new Thread("PRODUCER") {  
        public void run() {  
            String event = "FOUR";  
            try { queue.put(event); // thread will block here  
                System.out.printf("[%s] published event : %s %n",  
                    Thread.currentThread().getName(), event); }  
            catch (InterruptedException e) {  
                e.printStackTrace(); } } };  
    producer.start(); // starting publisher thread  
  
    Thread consumer = new Thread("CONSUMER") {  
        public void run() {  
            try { String event = queue.take(); // thread will block here  
                System.out.printf("[%s] consumed event : %s %n",  
                    Thread.currentThread().getName(), event); }  
            catch (InterruptedException e) { e.printStackTrace(); }  
        }  
    };  
    consumer.start(); // starting consumer thread  
}
```

# Sincronizare <-> comunicare

Programele comunică între ele nu numai pentru a-și comunica informații sub formă de mesaje ci și pentru a se sincroniza.

=>

Un semnal de sincronizare poate fi considerat și el că este un mesaj fără conținut ce se transmite între programe.

Cum se realizează acest lucru?



# Procese secvențiale comunicante. Rendez-vous simetric.

task-ul A: o comandă de emitere mesaj → SUSPENDARE ←  
task-ul B: o comandă de recepție de mesaj;

task-ul B: o comandă de recepție de mesaj → SUSPENDARE ←  
task-ul A: o comandă de emisie de mesaj;

task-uri sincronizate → datele (mesajul) sunt transferate

# Procese distribuite. Rendez-vous asimetric

- Comunicarea și sincronizarea între programele concurente se realizează în acest caz similar cu apelarea prin nume de către programul emițător a unei proceduri incluse în programul receptor,
  - lista cu parametrii asociați acestui apel fiind folosită ca un “canal” de comunicare pentru transmiterea de date între cele două programe.
- Doar programul apelant trebuie să cunoască numele programului apelat, nu și invers.

Exemple...