

# Linked Lists: Locking, Lock-Free, and Beyond ...

Companion slides for  
The Art of Multiprocessor  
Programming  
by Maurice Herlihy & Nir Shavit

# Coarse-Grained Synchronization

- Each method locks the object
  - Avoid contention using queue locks
  - Easy to reason about
    - In simple cases
  - Standard Java model
    - **Synchronized** blocks and methods
- So, are we done?

# Coarse-Grained Synchronization

- Sequential bottleneck
  - Threads “stand in line”
- Adding more threads
  - Does not improve throughput
  - Struggle to keep it from getting worse
- So why even use a multiprocessor?
  - Well, some apps inherently parallel ...

# First:

## Fine-Grained Synchronization

- Instead of using a single lock ..
- Split object into
  - Independently-synchronized components
- Methods conflict when they access
  - The same component ...
  - At the same time

# Set Interface

- Unordered collection of items
- No duplicates
- Methods
  - `add(x)` put x in set
  - `remove(x)` take x out of set
  - `contains(x)` tests if x in set

# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

**Add item to set**

# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

**Remove item from set**



# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

**Is item in set?**

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```



item of interest

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```




Usually hash code

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

Reference to next node

# The List-Based Set



Sorted with Sentinel nodes  
(min & max possible keys)

# Reasoning about Concurrent Objects

- Invariant
  - Property that always holds
- Established because
  - True when object is **created**
  - Truth **preserved** by each method
    - Each **step** of each method

# Specifically ...

- Invariants preserved by
  - `add()`
  - `remove()`
  - `contains()`
- Most steps are trivial
  - Usually one step tricky
  - Often linearization point



# Interference


- Invariants make sense only if
  - methods considered
  - are the only modifiers
- Language encapsulation helps
  - List nodes not visible outside class

# Interference

- Freedom from interference needed even for removed nodes
  - Some algorithms traverse removed nodes
  - Careful with **malloc()** & **free()**!
- Garbage-collection helps here

# Abstract Data Types

- Concrete representation



- Abstract Type
  - {a, b}

# Abstract Data Types

- Meaning of rep given by abstraction map

$$- S( \boxed{\text{ } \text{ } } \rightarrow \boxed{a \text{ } } \rightarrow \boxed{b \text{ } } \rightarrow \boxed{\text{ } \text{ } } ) = \{a, b\}$$

# Rep Invariant

- Which concrete values meaningful?
  - Sorted?
  - Duplicates?
- Rep invariant
  - Characterizes legal concrete reps
  - Preserved by methods
  - Relied on by methods

# Blame Game

- Rep invariant is a **contract**
- Suppose
  - **add()** leaves behind 2 copies of  $x$
  - **remove()** removes only 1
- Which one is incorrect?

# Blame Game

- Suppose
  - **add()** leaves behind 2 copies of  $x$
  - **remove()** removes only 1
- Which one is incorrect?
  - If rep invariant says *no duplicates*
    - **add()** is incorrect
  - Otherwise
    - **remove()** is incorrect

# Rep Invariant (partly)

- Sentinel nodes
  - tail reachable from head
- Sorted
- No duplicates



# Abstraction Map

- $S(\text{head}) =$ 
  - $\{ x \mid \text{there exists } a \text{ such that}$ 
    - $a \text{ reachable from head and}$
    - $a.\text{item} = x$
  - $\}$

# Sequential List Based Set

Add()




Remove()




# Sequential List Based Set


Add()




Remove()



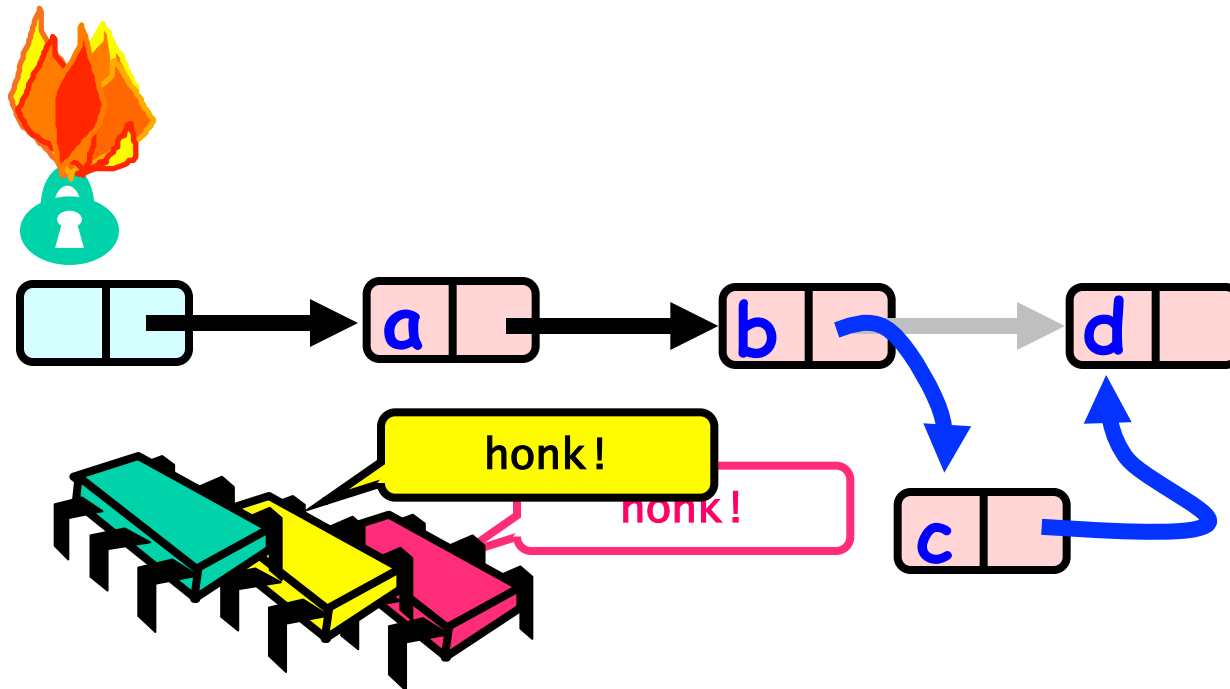
# Course Grained Locking



# Course Grained Locking



# Course Grained Locking



Simple but hotspot + bottleneck

# Coarse-Grained Locking



- Easy, same as synchronized methods
  - “One lock to rule them all ...”
- Simple, clearly correct
  - Deserves respect!
- Works poorly with contention
  - Queue locks help
  - But bottleneck still an issue

# Fine-grained Locking


- Requires careful thought
  - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”
- Split object into pieces
  - Each piece has own lock
  - Methods that work on disjoint pieces need not exclude each other




# Hand-over-Hand locking




# Hand-over-Hand locking




# Hand-over-Hand locking




# Hand-over-Hand locking




# Hand-over-Hand locking




# Removing a Node




remove(b)



# Removing a Node




# Removing a Node







# Removing a Node





# Removing a Node




# Removing a Node




# Concurrent Removes




# Concurrent Removes




# Concurrent Removes



# Concurrent Removes




# Concurrent Removes







# Concurrent Removes




# Concurrent Removes



# Concurrent Removes




# Uh, Oh




# Uh, Oh

Bad news, C not removed




# Problem

- To delete node *c*
  - Swing node *b*'s next field to *d*




- Problem is,
  - Someone deleting *b* concurrently could direct a pointer to *c*




# Insight

- If a node is locked
  - No one can delete node' s *successor*
- If a thread locks
  - Node to be deleted
  - And its predecessor
  - Then it works

# Hand-Over-Hand Again




remove(b)







# Hand-Over-Hand Again




# Hand-Over-Hand Again




# Hand-Over-Hand Again





# Hand-Over-Hand Again




# Hand-Over-Hand Again




# Removing a Node



# Removing a Node




# Removing a Node







# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node




# Removing a Node



# Removing a Node




# Removing a Node







# Removing a Node




# Removing a Node




# Removing a Node



# Removing a Node



# Removing a Node



# Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

# Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Key used to order node**

# Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        currNode.unlock();  
        predNode.unlock();  
    }  
}
```

**Predecessor and current nodes**



# Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;
```

```
    try {
```

```
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Make sure  
locks released**

# Remove method

```
public boolean remove(Item item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Everything else**


# Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```

# Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```


**lock pred == head**



# Remove method

```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```


Lock current



# Remove method


```
try {  
    pred = this.head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    ...  
} finally { ... }
```

Traversing list



# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



# Remove: searching


```
while (curr.key <= key) {
```

```
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }
```

```
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

```
return false;
```

Search key range






# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```


**At start of each loop: curr  
and pred locked**



# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```


**If item found, remove node**



## Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```


## If node found, remove it



# Remove: searching

Unlock predecessor

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```




# Remove: searching

**Only one node locked!**

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```


**pred.unlock();**



# Remove: searching


```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

demote current



# Remove: searching


```
while (curr.key <= key) {  
    Find and lock new current  
    if (curr == curr.next) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = currNode;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = currNode;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

Lock invariant restored





# Remove: searching

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

**Otherwise, not present**



**return false;**

# Why does this work?

- To remove node  $e$ 
  - Must lock  $e$
  - Must lock  $e$ 's predecessor
- Therefore, if you lock a node
  - It can't be removed
  - And neither can its successor

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

- pred reachable from head
- curr is pred.next
- So curr.item is in the set

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point if  
item is present**

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

Node locked, so no other  
thread can remove it ....

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

**return false;**

Item not present

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

**return false;**

- pred reachable from head
- curr is pred.next
- pred.key < key
- key < curr.key

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point**





# Adding Nodes

- To add node  $e$ 
  - Must lock predecessor
  - Must lock successor
- Neither can be deleted
  - (Is successor lock actually required?)

# Same Abstraction Map

- $S(\text{head}) =$ 
  - $\{ x \mid \text{there exists } a \text{ such that}$ 
    - $a \text{ reachable from head and}$
    - $a.\text{item} = x$
  - $\}$

# Rep Invariant

- Easy to check that
  - tail always reachable from head
  - Nodes sorted, no duplicates

# Drawbacks

- Better than coarse-grained lock
  - Threads can traverse in parallel
- Still not ideal
  - Long chain of acquire/release
  - Inefficient

# This work is licensed under a

## [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/3.0/).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.